

Maze Shaping Report: Manhattan PBRs with Revisit-Termination (v2 Long-Train)

February 13, 2026

1 Setup

This run keeps the same revisit-termination rule but increases the training budget from 500 episodes to 5000 episodes.

Core settings:

- Termination: goal reached or revisit of any previously visited state in same episode
- Potential: Manhattan PBRs
- Discount: $\gamma = 0.99$
- Rewards: step = 0.0, goal = 1.0
- Episodes: 5000, Runs: 12
- Validation: every 25 episodes (plus episode 1), 30 rollouts

2 Curves




Figure 1: Training curve (steps to termination).




Figure 2: Validation success rate and validation average steps to termination.

Condition	Train mean steps	Validation success rate	Validation mean steps
no_shaping	1.916667	0.0	1.000000
phi_half	9.166667	0.0	10.083333
phi_full	5.333333	0.0	8.416667

Table 1: Episode-5000 summary under revisit-termination rule.

3 Final Metrics (Episode 5000)

4 Interpretation

Increasing training iterations alone did not recover success under this termination rule (all conditions remained at 0 validation success). This supports the view that the termination design, not only training budget, is the primary bottleneck.

5 Artifacts

- Output root: `.../outputs/maze_shaping_revisit_termination_manhattan_v2_longtrain`
- Files: `learning_curve.csv/png`, `validation_progress.csv/png`, `run_summary.json`